

DATA SCIENCE

Building a Modern Dashboard with Python and Gradio

Data insights made simple

Thomas Reid

Jun 4, 2025 13 min read



Image by AI (Gpt4-o)

This is the second in a short series on developing data dashboards using the latest Python-based GUI development tools, Streamlit, Gradio, and Taipy.

The source dataset for each dashboard will be the same, but stored in different formats. As much as possible, I'll also try to make the actual dashboard layouts for each tool resemble each other and have the same functionality.

In the first part of this series, I created a Streamlit version of the dashboard that retrieves its data from a local PostgreSQL database. You can view that article [here](#).

This time, we're exploring the use of the Gradio library.

The data for this dashboard will be in a local CSV file, and Pandas will be our primary data processing engine.

If you want to see a quick demo of the app, I have deployed it to Hugging Face Spaces. You can run it using the link below, but note that the two input date picker pop-ups do not work due to a known bug in the Hugging Face environment. This is only the case for deployed apps on HF, you can still change the dates manually. Running the app locally works fine and doesn't have this issue.

Dashboard demo on HuggingFace

What is Gradio?

Gradio is an open-source Python package that simplifies the process of building demos or web applications for machine learning models, APIs, or any Python function. With it, you can create demos or web applications without needing JavaScript, CSS, or web hosting experience. By writing just a few lines of Python code, you can unlock the power of Gradio and seamlessly showcase your machine-learning models to a broader audience.

Gradio simplifies the development process by providing an intuitive framework that eliminates the complexities associated with building user interfaces from scratch. Whether you are a machine learning developer, researcher, or enthusiast, Gradio allows you to create beautiful and

interactive demos that enhance the understanding and accessibility of your machine learning models.

This open-source Python package helps you bridge the gap between your machine learning expertise and a broader audience, making your models accessible and actionable.

What we'll develop

We're developing a data dashboard. Our source data will be a single CSV file containing 100,000 synthetic sales records.

The actual source of the data isn't *that* important. It could just as easily be a text file, an Excel file, SQLite, or any database you can connect to.

This is what our final dashboard will look like.

Sales Performance Dashboard


Start Date 2010-01-01	End Date 2023-12-30	Category All Categories
--------------------------	------------------------	----------------------------

Key Metrics

Total Revenue 28140530.550100002	Total Orders 100000	Average Order Value 281.4053055	Top Category Electronics
-------------------------------------	------------------------	------------------------------------	-----------------------------

Visualisations

Revenue Over Time Revenue by Category Top Products



Raw Data

order_id	order_date	customer_id	customer_name	product_id	product_name	categories	quantity	price	total	revenue
0	2022-08-01 00:00:00	245	Customer_884	281	Smartphone	Electronics	3	98.82	270.06	270.06
1	2022-02-19 00:00:00	701	Customer_1672	285	Printer	Electronics	6	12.74	76.44	76.44
2	2017-01-01 00:00:00	184	Customer_21728	288	Notebook	Stationery	8	48.35	386.8	386.8
3	2013-03-09 00:00:00	275	Customer_23778	288	Laptop	Electronics	3	74.85	224.55	224.549999999
4	2022-04-23 00:00:00	960	Customer_23798	218	Cabinet	Office	6	53.77	322.62	322.62
5	2019-07-18 00:00:00	197	Customer_25587	202	Desk	Office	3	47.17	141.51	141.51
6	2014-11-12 00:00:00	510	Customer_6912	204	Monitor	Electronics	5	22.5	112.5	112.5
7	2016-07-12 00:00:00	150	Customer_17761	208	Laptop	Electronics	9	49.33	443.97	443.969999999
8	2016-11-12 00:00:00	997	Customer_23801	209	Coffee Maker	Electronics	7	47.22	330.54	330.539999999
9	2017-01-23 00:00:00	151	Customer_38325	207	Pan	Stationery	6	3.5	21	21
10	2013-11-27 00:00:00	886	Customer_14540	206	Paper	Stationery	8	91.63	733.04	733.04

Image by Author

There are four main sections.

- The top row enables the user to select specific start and end dates and/or product categories using date pickers and a drop-down list, respectively.

- The second row—Key metrics—shows a top-level summary of the chosen data.
- The Visualisation section allows the user to select one of three graphs to display the input dataset.
- The raw data section is precisely what it claims to be. This tabular representation of the chosen data effectively shows a snapshot of the underlying CSV data file.

Using the dashboard is easy. Initially, stats for the whole data set are displayed. The user can then narrow the data focus using the three filter fields at the top of the display. The graphs, key metrics, and raw data sections dynamically update to reflect the user's choices in the filter fields.

The underlying data

As mentioned, the dashboard's source data is contained in a single comma-separated values (CSV) file. The data consists of 100,000 synthetic sales-related records. Here are the first ten records of the file to give you an idea of what it looks like.

order_id	order_date	customer_id	customer_name	product_id	product
0	01/08/2022	245	Customer_884	201	Smartph
1	19/02/2022	701	Customer_1672	205	Printer
2	01/01/2017	184	Customer_21720	208	Notebo
3	09/03/2013	275	Customer_23770	200	Laptop
4	23/04/2022	960	Customer_23790	210	Cabinet
5	10/07/2019	197	Customer_25587	202	Desk
6	12/11/2014	510	Customer_6912	204	Monitor
7	12/07/2016	150	Customer_17761	200	Laptop
8	12/11/2016	997	Customer_23801	209	Coffee

9	23/01/2017	151	Customer_30325	207	Pen
+-----+	+-----+	+-----+	+-----+	+-----+	+-----+

And here is some Python code you can use to generate a similar dataset. Ensure that both the NumPy and Pandas libraries are installed first.

```
# generate the 100K record CSV file
#
import polars as pl
import numpy as np
from datetime import datetime, timedelta

def generate(nrows: int, filename: str):
    names = np.asarray(
        [
            "Laptop",
            "Smartphone",
            "Desk",
            "Chair",
            "Monitor",
            "Printer",
            "Paper",
            "Pen",
            "Notebook",
            "Coffee Maker",
            "Cabinet",
            "Plastic Cups",
        ]
    )
    categories = np.asarray(
        [
            "Electronics",
            "Electronics",
            "Office",
            "Office",
            "Electronics",
            "Electronics",
            "Stationery",
        ]
    )
```

```

        "Stationery",
        "Stationery",
        "Electronics",
        "Office",
        "Sundry",
    ]
)
product_id = np.random.randint(len(names), size=nrows)
quantity = np.random.randint(1, 11, size=nrows)
price = np.random.randint(199, 10000, size=nrows) / 100
# Generate random dates between 2010-01-01 and 2023-12-31
start_date = datetime(2010, 1, 1)
end_date = datetime(2023, 12, 31)
date_range = (end_date - start_date).days
# Create random dates as np.array and convert to string format
order_dates = np.array([(start_date + timedelta(days=np.random.randint(0, date_range)))
# Define columns
columns = {
    "order_id": np.arange(nrows),
    "order_date": order_dates,
    "customer_id": np.random.randint(100, 1000, size=nrows),
    "customer_name": [f"Customer_{i}" for i in np.random.randint(2**15,
    "product_id": product_id + 200,
    "product_names": names[product_id],
    "categories": categories[product_id],
    "quantity": quantity,
    "price": price,
    "total": price * quantity,
}
# Create Polars DataFrame and write to CSV with explicit delimiter
df = pl.DataFrame(columns)
df.write_csv(filename, separator=',', include_header=True) # Ensure comma

# Generate 100,000 rows of data with random order_date and save to CSV
generate(100_000, "/mnt/d/sales_data/sales_data.csv")

```

Installing and using Gradio

Installing Gradio is easy using **pip**, but for coding, the best practice is to set up a separate Python environment for all your work. I use Miniconda for that purpose, but feel free to use whatever method suits your work practice.

If you want to go down the conda route and don't already have it, you must install Miniconda (recommended) or Anaconda first.

Please note that, at the time of writing, Gradio needs at least Python 3.8 installed to work correctly.

Once the environment is created, switch to it using the **'activate'** command, and then run **'pip install'** to install our required Python libraries.

```
#create our test environment
(base) C:\Users\thoma>conda create -n gradio_dashboard python=3.12 -y

# Now activate it
(base) C:\Users\thoma>conda activate gradio_dashboard

# Install python libraries, etc ...
(gradio_dashboard) C:\Users\thoma>pip install gradio pandas matplotlib cache
```

Key differences between Streamlit and Gradio

As I'll demonstrate in this article, it's possible to produce very similar data dashboards using Streamlit and Gradio. However, their ethos differs in several key ways.

Focus

- Gradio specialises in creating interfaces for machine learning models, whilst Streamlit is more designed for general-purpose data applications and visualisations.

Ease of use

- Gradio is known for its simplicity and rapid prototyping capabilities, making it easier for beginners to use. Streamlit offers more advanced features and customisation options, which may require a steeper learning curve.

Interactivity

- Streamlit uses a reactive programming model where any input change triggers a complete script rerun, updating all components immediately. Gradio, by default, updates only when a user clicks a submit button, though it can be configured for live updates.

Customization

- Gradio focuses on pre-built components for quickly demonstrating AI models. Streamlit provides more extensive customisation options and flexibility for complex projects.

Deployment

- Having deployed both a Streamlit and a Gradio app, I would say it's easier to deploy a Streamlit app than a Gradio app. In Streamlit, deployment can be done with a single click via the Streamlit Community Cloud. This functionality is built into any Streamlit app you create. Gradio offers deployment using Hugging Face Spaces, but it involves more work. Neither method is particularly complex, though.

Use cases

Streamlit excels in creating data-centric applications and interactive dashboards for complex projects. Gradio is ideal for quickly showcasing machine learning models and building simpler applications.

The Gradio Dashboard Code

I'll break down the code into sections and explain each one as we proceed.

We begin by importing the required external libraries and loading the full dataset from the CSV file into a Pandas DataFrame.

```
import gradio as gr
import pandas as pd
import matplotlib.pyplot as plt
import datetime
import warnings
import os
import tempfile
from cachetools import cached, TTLCache

warnings.filterwarnings("ignore", category=FutureWarning, module="seaborn")

# -----
# 1) Load CSV data once
# -----
csv_data = None

def load_csv_data():
    global csv_data

    # Optional: specify column dtypes if known; adjust as necessary
    dtype_dict = {
        "order_id": "Int64",
        "customer_id": "Int64",
        "product_id": "Int64",
        "quantity": "Int64",
        "price": "float",
```

```
"total": "float",
"customer_name": "string",
"product_names": "string",
"categories": "string"
}

csv_data = pd.read_csv(
    "d:/sales_data/sales_data.csv",
    parse_dates=["order_date"],
    dayfirst=True,          # if your dates are DD/MM/YYYY format
    low_memory=False,
    dtype=dtype_dict
)

load_csv_data()
```

Next, we configure a time-to-live cache with a maximum of 128 items and an expiration of 300 seconds. This is used to store the results of expensive function calls and speed up repeated lookups

The **get_unique_categories** function returns a list of unique, cleaned (capitalised) categories from the `csv_data` DataFrame, caching the result for quicker access.

The **get_date_range** function returns the minimum and maximum order dates from the dataset, or None if the data is unavailable.

The **filter_data** function filters the `csv_data` DataFrame based on a specified date range and optional category, returning the filtered DataFrame.

The **get_dashboard_stats** function retrieves summary metrics—total revenue, total orders, average order value, and top category—for the given filters. Internally it uses `filter_data()` to scope the dataset and then calculate these key statistics.

The **get_data_for_table** function returns a detailed DataFrame of filtered sales data, sorted by **order_id** and **order_date**, including additional revenue for each sale.

The **get_plot_data** function formats data for generating a plot by summing revenue over time, grouped by date.

The **get_revenue_by_category** function aggregates and returns revenue by category, sorted by revenue, within the specified date range and category.

The **get_top_products** function returns the top 10 products by revenue, filtered by date range and category.

Based on the orientation argument, the **create_matplotlib_figure** function generates a bar plot from the data and saves it as an image file, either vertical or horizontal.

```
cache = TTLCache(maxsize=128, ttl=300)

@cached(cache)
def get_unique_categories():
    global csv_data
    if csv_data is None:
        return []
    cats = sorted(csv_data['categories'].dropna().unique().tolist())
    cats = [cat.capitalize() for cat in cats]
    return cats

def get_date_range():
    global csv_data
    if csv_data is None or csv_data.empty:
        return None, None
    return csv_data['order_date'].min(), csv_data['order_date'].max()

def filter_data(start_date, end_date, category):
    global csv_data
```

```

if isinstance(start_date, str):
    start_date = datetime.datetime.strptime(start_date, '%Y-%m-%d').date()
if isinstance(end_date, str):
    end_date = datetime.datetime.strptime(end_date, '%Y-%m-%d').date()

df = csv_data.loc[
    (csv_data['order_date'] >= pd.to_datetime(start_date)) &
    (csv_data['order_date'] <= pd.to_datetime(end_date))
].copy()

if category != "All Categories":
    df = df.loc[df['categories'].str.capitalize() == category].copy()

return df

def get_dashboard_stats(start_date, end_date, category):
    df = filter_data(start_date, end_date, category)
    if df.empty:
        return (0, 0, 0, "N/A")

    df['revenue'] = df['price'] * df['quantity']
    total_revenue = df['revenue'].sum()
    total_orders = df['order_id'].nunique()
    avg_order_value = total_revenue / total_orders if total_orders else 0

    cat_revenues = df.groupby('categories')['revenue'].sum().sort_values(ascending=True)
    top_category = cat_revenues.index[0] if not cat_revenues.empty else "N/A"

    return (total_revenue, total_orders, avg_order_value, top_category.capitalize())

def get_data_for_table(start_date, end_date, category):
    df = filter_data(start_date, end_date, category)
    if df.empty:
        return pd.DataFrame()

    df = df.sort_values(by=["order_id", "order_date"], ascending=[True, False])

    columns_order = [
        "order_id", "order_date", "customer_id", "customer_name",

```

```

        "product_id", "product_names", "categories", "quantity",
        "price", "total"
    ]
    columns_order = [col for col in columns_order if col in df.columns]
    df = df[columns_order].copy()

    df['revenue'] = df['price'] * df['quantity']
    return df

def get_plot_data(start_date, end_date, category):
    df = filter_data(start_date, end_date, category)
    if df.empty:
        return pd.DataFrame()
    df['revenue'] = df['price'] * df['quantity']
    plot_data = df.groupby(df['order_date'].dt.date)['revenue'].sum().reset_index()
    plot_data.rename(columns={'order_date': 'date'}, inplace=True)
    return plot_data

def get_revenue_by_category(start_date, end_date, category):
    df = filter_data(start_date, end_date, category)
    if df.empty:
        return pd.DataFrame()
    df['revenue'] = df['price'] * df['quantity']
    cat_data = df.groupby('categories')['revenue'].sum().reset_index()
    cat_data = cat_data.sort_values(by='revenue', ascending=False)
    return cat_data

def get_top_products(start_date, end_date, category):
    df = filter_data(start_date, end_date, category)
    if df.empty:
        return pd.DataFrame()
    df['revenue'] = df['price'] * df['quantity']
    prod_data = df.groupby('product_names')['revenue'].sum().reset_index()
    prod_data = prod_data.sort_values(by='revenue', ascending=False).head(10)
    return prod_data

def create_matplotlib_figure(data, x_col, y_col, title, xlabel, ylabel, orient):
    plt.figure(figsize=(10, 6))
    if data.empty:
        plt.text(0.5, 0.5, 'No data available', ha='center', va='center')

```

```

else:
    if orientation == 'v':
        plt.bar(data[x_col], data[y_col])
        plt.xticks(rotation=45, ha='right')
    else:
        plt.barh(data[x_col], data[y_col])
        plt.gca().invert_yaxis()

plt.title(title)
plt.xlabel(xlabel)
plt.ylabel(ylablel)
plt.tight_layout()

with tempfile.NamedTemporaryFile(delete=False, suffix=".png") as tmpfile:
    plt.savefig(tmpfile.name)
plt.close()
return tmpfile.name

```

The **update_dashboard** function retrieves key sales statistics (total revenue, total orders, average order value, and top category) by calling the `get_dashboard_stats` function. It gathers data for three distinct visualisations (revenue over time, revenue by category, and top products), then uses `create_matplotlib_figure` to generate plots. It prepares and returns a data table (via the `get_data_for_table()` function) along with all generated plots and stats so they can be displayed in the dashboard.

The **create_dashboard** function sets the date boundaries (minimum and maximum dates) and establishes the initial default filter values. It uses Gradio to construct a user interface (UI) featuring date pickers, category drop-downs, key metric displays, plot tabs, and a data table. It then wires up the filters so that changing any of them triggers a call to the **update_dashboard** function, ensuring the dashboard visuals and metrics are always in sync with the selected filters. Finally, it returns the assembled Gradio interface launched as a web application.

```

def update_dashboard(start_date, end_date, category):
    total_revenue, total_orders, avg_order_value, top_category = get_dashboard_data(start_date, end_date, category)

    # Generate plots
    revenue_data = get_plot_data(start_date, end_date, category)
    category_data = get_revenue_by_category(start_date, end_date, category)
    top_products_data = get_top_products(start_date, end_date, category)

    revenue_over_time_path = create_matplotlib_figure(
        revenue_data, 'date', 'revenue',
        "Revenue Over Time", "Date", "Revenue"
    )
    revenue_by_category_path = create_matplotlib_figure(
        category_data, 'categories', 'revenue',
        "Revenue by Category", "Category", "Revenue"
    )
    top_products_path = create_matplotlib_figure(
        top_products_data, 'product_names', 'revenue',
        "Top Products", "Revenue", "Product Name", orientation='h'
    )

    # Data table
    table_data = get_data_for_table(start_date, end_date, category)

    return (
        revenue_over_time_path,
        revenue_by_category_path,
        top_products_path,
        table_data,
        total_revenue,
        total_orders,
        avg_order_value,
        top_category
    )

def create_dashboard():
    min_date, max_date = get_date_range()
    if min_date is None or max_date is None:
        min_date = datetime.datetime.now()
        max_date = datetime.datetime.now()

```



```

default_start_date = min_date
default_end_date = max_date

with gr.Blocks(css="""
    footer {display: none !important;}
    .tabs {border: none !important;}
    .gr-plot {border: none !important; box-shadow: none !important;}
""") as dashboard:

    gr.Markdown("# Sales Performance Dashboard")

    # Filters row
    with gr.Row():
        start_date = gr.DateTime(
            label="Start Date",
            value=default_start_date.strftime('%Y-%m-%d'),
            include_time=False,
            type="datetime"
        )
        end_date = gr.DateTime(
            label="End Date",
            value=default_end_date.strftime('%Y-%m-%d'),
            include_time=False,
            type="datetime"
        )
        category_filter = gr.Dropdown(
            choices=["All Categories"] + get_unique_categories(),
            label="Category",
            value="All Categories"
        )

    gr.Markdown("# Key Metrics")

    # Stats row
    with gr.Row():
        total_revenue = gr.Number(label="Total Revenue", value=0)
        total_orders = gr.Number(label="Total Orders", value=0)
        avg_order_value = gr.Number(label="Average Order Value", value=0)
        top_category = gr.Textbox(label="Top Category", value="N/A")

```

```

gr.Markdown("# Visualisations")
# Tabs for Plots
with gr.Tabs():
    with gr.Tab("Revenue Over Time"):
        revenue_over_time_image = gr.Image(label="Revenue Over Time")
    with gr.Tab("Revenue by Category"):
        revenue_by_category_image = gr.Image(label="Revenue by Category")
    with gr.Tab("Top Products"):
        top_products_image = gr.Image(label="Top Products", container_width=300)

gr.Markdown("# Raw Data")
# Data Table (below the plots)
data_table = gr.DataFrame(
    label="Sales Data",
    type="pandas",
    interactive=False
)

# When filters change, update everything
for f in [start_date, end_date, category_filter]:
    f.change(
        fn=lambda s, e, c: update_dashboard(s, e, c),
        inputs=[start_date, end_date, category_filter],
        outputs=[
            revenue_over_time_image,
            revenue_by_category_image,
            top_products_image,
            data_table,
            total_revenue,
            total_orders,
            avg_order_value,
            top_category
        ]
    )

# Initial load
dashboard.load(
    fn=lambda: update_dashboard(default_start_date, default_end_date, default_category),
    outputs=[

```

```
        revenue_over_time_image,  
        revenue_by_category_image,  
        top_products_image,  
        data_table,  
        total_revenue,  
        total_orders,  
        avg_order_value,  
        top_category  
    ]  
)  
  
    return dashboard  
  
if __name__ == "__main__":  
    dashboard = create_dashboard()  
    dashboard.launch(share=False)
```

Running the program

Create a Python file, e.g. `gradio_test.py`, and insert all the above code snippets. Save it, and run it like this,

```
(gradio_dashboard) $ python gradio_test.py
```

* Running on local URL: `http://127.0.0.1:7860`

To create a public link, `set `share=True` in `launch()``.

Click on the local URL shown, and the dashboard will open full screen in your browser.

Summary

This article provides a comprehensive guide to building an interactive sales performance dashboard using Gradio and a CSV file as its source data.

Gradio is a modern, Python-based open-source framework that simplifies the creation of data-driven dashboards and GUI applications. The dashboard I developed allows users to filter data by date ranges and product categories, view key metrics such as total revenue and top-performing categories, explore visualisations like revenue trends and top products, and navigate through raw data with pagination.

I also mentioned some key differences between developing visualisation tools using Gradio and Streamlit, another popular front-end Python library.

This guide provides a comprehensive implementation of a Gradio data dashboard, covering the entire process from creating sample data to developing Python functions for querying data, generating plots, and handling user input. This step-by-step approach demonstrates how to leverage Gradio's capabilities to create user-friendly and dynamic dashboards, making it ideal for data engineers and scientists who want to build interactive data applications.

Although I used a CSV file for my data, modifying the code to use another data source, such as a relational database management system (RDBMS) like SQLite, should be straightforward. For example, in my other article in this series on creating a similar dashboard using Streamlit, the data source is a PostgreSQL database.